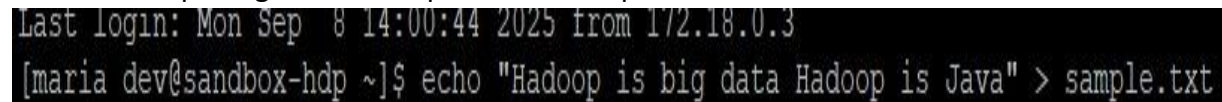


## PRACTICAL-04

**AIM:-** Hadoop WordCount Execution Steps (Using PuTTY).

### **Step 1:** Prepare Input Data

```
echo "Hadoop is big data Hadoop is Java" > sample.txt
```

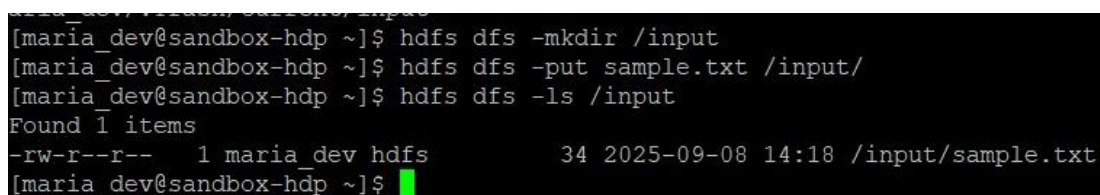


```
Last login: Mon Sep  8 14:00:44 2025 from 172.18.0.3  
[maria_dev@sandbox-hdp ~]$ echo "Hadoop is big data Hadoop is Java" > sample.txt
```

```
hdfs dfs -mkdir /input
```

```
hdfs dfs -put sample.txt /input/
```

```
hdfs dfs -ls /input
```



```
[maria_dev@sandbox-hdp ~]$ hdfs dfs -mkdir /input  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -put sample.txt /input/  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -ls /input  
Found 1 items  
-rw-r--r--  1 maria_dev hdfs          34 2025-09-08 14:18 /input/sample.txt  
[maria_dev@sandbox-hdp ~]$
```

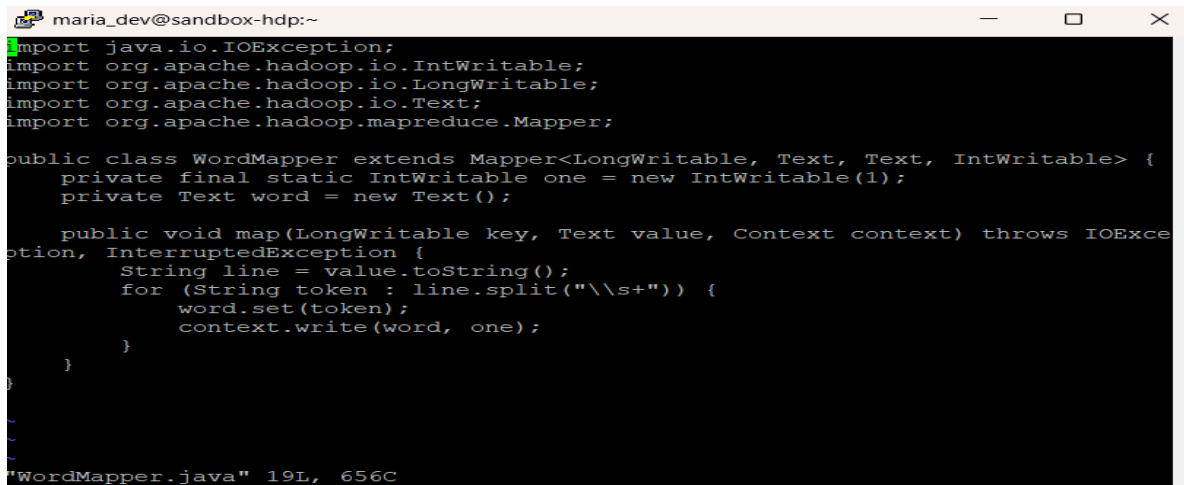
### **Step 2:** Create Java Files

**By using this code we will create the java files by using vi command and then pasting the above code and then save the file.**

#### **WordMapper.java**

```
import java.io.IOException;  
import  
org.apache.hadoop.io.IntWritable;  
import  
org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
    public void map(LongWritable key, Text value, Context context) throws IOException,  
InterruptedException {  
        String line = value.toString();
```

```
        for (String token : line.split("\\s+")) { // Corrected escape sequence
            word.set(token);
            context.write(word, one);
        }
    }
}
```

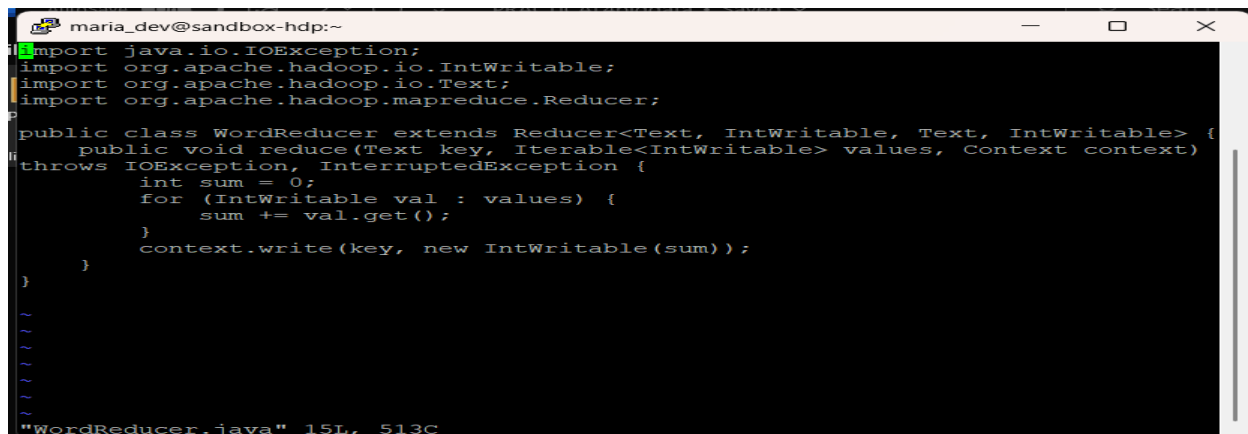
A screenshot of a terminal window titled 'maria\_dev@sandbox-hdp:~'. The terminal displays the Java code for 'WordMapper.java'. The code includes imports for 'java.io.IOException', 'org.apache.hadoop.io.IntWritable', 'org.apache.hadoop.io.LongWritable', 'org.apache.hadoop.io.Text', and 'org.apache.hadoop.mapreduce.Mapper'. It defines a 'WordMapper' class that extends 'Mapper<LongWritable, Text, Text, IntWritable>'. The class has a 'map' method that takes 'LongWritable key', 'Text value', and 'Context context' as parameters. Inside the 'map' method, it converts the 'value' to a 'String line', splits it by whitespace using 'line.split("\\s+")', and iterates over the tokens to set a 'Text word' and write it to the 'context' with a value of 'one'. The terminal shows the code from line 19 to line 656.

```
maria_dev@sandbox-hdp:~  
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
  
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
  
    public void map(LongWritable key, Text value, Context context) throws IOExce  
ption, InterruptedException {  
        String line = value.toString();  
        for (String token : line.split("\\s+")) {  
            word.set(token);  
            context.write(word, one);  
        }  
    }  
}
```

"WordMapper.java" 19L, 656C

## WordReducer.java

```
import java.io.IOException;  
import  
org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;  
public class WordReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws  
IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val :  
            values) { sum +=  
                val.get();  
            }  
  
        context.write(key, new IntWritable(sum));  
    }  
}
```



```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class WordReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

"WordReducer.java" 15L, 513C

## WordCountDriver.java

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WordCountDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: WordCountDriver <input path> <output path>");
            System.exit(-1);
        }

        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCountDriver.class);
        job.setMapperClass(WordMapper.class);
        job.setReducerClass(WordReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

```

maria_dev@sandbox-hdp:~
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCountDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: WordCountDriver <input path> <output path>");
            System.exit(-1);
        }
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCountDriver.class);
        job.setMapperClass(WordMapper.class);
        job.setReducerClass(WordReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
    }
}

```

```

[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java
[maria_dev@sandbox-hdp ~]$ vi WordCountDriver.java
[maria_dev@sandbox-hdp ~]$ mkdir wordcount_classes
[maria_dev@sandbox-hdp ~]$ javac -cp 'hadoop classpath' -d wordcount_classes WordMapper.java WordReducer.java WordCountDriver.java

```

## Step 3: Compile Java Files

mkdir wordcount\_classes

```

[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java
[maria_dev@sandbox-hdp ~]$ vi WordCountDriver.java
[maria_dev@sandbox-hdp ~]$ mkdir wordcount_classes
[maria_dev@sandbox-hdp ~]$ javac -cp 'hadoop classpath' -d wordcount_classes WordMapper.java WordReducer.java WordCountDriver.java

```

hadoop com.sun.tools.javac.Main -d wordcount\_classes WordMapper.java WordReducer.java WordCountDriver.java

(If error, use:)

javac -cp `hadoop classpath` -d wordcount\_classes WordMapper.java WordReducer.java WordCountDriver.java

## Step 4: Create JAR

jar -cvf wordcount.jar -C wordcount\_classes/ .

```

[maria_dev@sandbox-hdp ~]$ jar -cvf wordcount.jar -C wordcount_classes/ .
added manifest
adding: WordMapper.class(in = 1867) (out= 776) (deflated 58%)
adding: WordReducer.class(in = 1592) (out= 665) (deflated 58%)
adding: WordCountDriver.class(in = 1535) (out= 853) (deflated 44%)
[maria_dev@sandbox-hdp ~]$

```

## Step 5: Run MapReduce Job

```
hdfs dfs -rm -r /output
```

```
hadoop jar wordcount.jar WordCountDriver /input /output
```

```
[maria_dev@sandbox-hdp ~]$ hadoop jar wordcount.jar WordCountDriver /input /output
25/09/09 06:15:06 INFO client.RMPProxy: Connecting to ResourceManager at sandbox-hdp.hortonworks.com/172.18.0.2:8032
25/09/09 06:15:06 INFO client.AHSProxy: Connecting to Application History server at sandbox-hdp.hortonworks.com/172.18.0.2:10200
25/09/09 06:15:08 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
25/09/09 06:15:10 INFO input.FileInputFormat: Total input paths to process : 1
25/09/09 06:15:11 INFO mapreduce.JobSubmitter: number of splits:1
25/09/09 06:15:12 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1757396666364_0001
25/09/09 06:15:14 INFO impl.YarnClientImpl: Submitted application application_1757396666364_0001
25/09/09 06:15:15 INFO mapreduce.Job: The url to track the job: http://sandbox-hdp.hortonworks.com:8088/proxy/application_1757396666364_0001/
25/09/09 06:15:15 INFO mapreduce.Job: Running job: job_1757396666364_0001
25/09/09 06:18:29 INFO mapreduce.Job: Job job_1757396666364_0001 running in uber mode : false
25/09/09 06:18:29 INFO mapreduce.Job: map 0% reduce 0%
25/09/09 06:20:12 INFO mapreduce.Job: map 100% reduce 0%
25/09/09 06:20:52 INFO mapreduce.Job: map 100% reduce 100%
25/09/09 06:20:58 INFO mapreduce.Job: Job job_1757396666364_0001 completed successfully
25/09/09 06:20:59 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=82
    FILE: Number of bytes written=305979
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=155
    HDFS: Number of bytes written=34
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=91699
    Total time spent by all reduces in occupied slots (ms)=37770
    Total time spent by all map tasks (ms)=91699
    Total time spent by all reduce tasks (ms)=37770
    Total vcore-milliseconds taken by all map tasks=91699
    Total vcore-milliseconds taken by all reduce tasks=37770
    Total megabyte-milliseconds taken by all map tasks=22924750
    Total megabyte-milliseconds taken by all reduce tasks=9442500
  Map-Reduce Framework
    Map input records=1
    Map output records=7
    Map output bytes=62
    Map output materialized bytes=82
```

## Step 6: View Output

```
hdfs dfs -ls /output
```

```
[maria_dev@sandbox-hdp ~]$ hdfs dfs -ls /output
Found 2 items
-rw-r--r--  1 maria_dev hdfs          0 2025-09-09 06:20 /output/_SUCCESS
-rw-r--r--  1 maria_dev hdfs        34 2025-09-09 06:20 /output/part-r-00000
```

```
hdfs dfs -cat /output/part-r-00000
```

```
[maria_dev@sandbox-hdp ~]$ hdfs dfs -cat /output/part-r-00000
Hadoop 2
Java 1
big 1
data 1
is 2
[maria_dev@sandbox-hdp ~]$
```

**CONCLUSION:-** We have completed the Hadoop WordCount Execution steps using Putty.